

NCBITaxonomy.jl - rapid biological names finding and reconciliation

Timothée Poisot^{1,2} Rory Gibb^{3,4} Sadie J. Ryan^{5,6,7} Colin J. Carlson^{8,9}

¹ Université de Montréal, Département de Sciences Biologiques, Montréal QC, Canada ² Québec Centre for Biodiversity Science, Montréal, QC, Canada ³ Centre on Climate Change and Planetary Health, London School of Hygiene and Tropical Medicine, London, UK ⁴ Centre for Mathematical Modelling of Infectious Diseases, London School of Hygiene and Tropical Medicine, London, UK ⁵ Emerging Pathogens Institute, University of Florida, Gainesville, FL, United States of America ⁶ School of Life Sciences, University of KwaZulu-Natal, Durban, South Africa ⁷ Department of Geography, University of Florida, Gainesville, FL, United States of America ⁸ Department of Microbiology and Immunology, Georgetown University Medical Center, Georgetown University, Washington, D.C., United States of America ⁹ Center for Global Health Science and Security, Georgetown University Medical Center, Georgetown University, Washington, D.C., United States of America

Correspondance to:

Timothée Poisot — timothee.poisot@umontreal.ca

This work is released by its authors under a CC-BY 4.0 license



Last revision: *October 14, 2021*

NCBITaxonomy.jl is a package designed to facilitate the reconciliation and cleaning of taxonomic names, using a local copy of the NCBI taxonomic backbone (Federhen 2012, Schoch et al. 2020); The basic search functions are coupled with quality-of-life functions including case-insensitive search and custom fuzzy string matching to facilitate the amount of information that can be extracted automatically while allowing efficient manual curation and inspection of results. NCBITaxonomy.jl works with version 1.6 of the Julia programming language (Bezanson et al. 2017), and relies on the Apache Arrow format to store a local copy of the NCBI raw taxonomy files. The design of NCBITaxonomy.jl has been inspired by similar efforts, like the R package taxadb (Norman et al. 2020), which provides an offline alternative to packages like taxize (Chamberlain and Szöcs 2013).

1 Unambiguously identifying species is a far more challenging task than it may appear. There are a vast
2 number of reasons for this. Different databases keep different taxonomic “backbones,” *i.e.* different data
3 structures in which names are mapped to species, and organised in a hierarchy. Not all names are unique
4 identifiers to groups. For example, *Io* can either refer to a genus of plants from the aster family, or to a
5 genus of molluscs; the genus *Mus* (of which the house mouse *Mus musculus* is a species), contains a
6 sub-genus *also* named *Mus*. Conversely, the same species can have several names, which are valid
7 synonyms: for example, the domestic cow *Bos taurus* admits *Bos primigenius taurus* as a valid synonym.
8 Taxonomic nomenclature also changes regularly, with groups being split, merged, or moved to a new
9 position in the tree of life; this is, notably, a common occurrence with viral taxonomy, each subsequent
10 version of which can differ markedly from the last; compare, *e.g.* Lefkowitz et al. (2018) to Walker et al.
11 (2020).

12 To add to the complexity, one must also consider that most taxa names are at some point manually typed,
13 which has the potential to introduce additional mistakes in raw data; it is likely to expect that such
14 mistakes may arise when attempting to write down the (perfectly valid) names of the bacterial isolate
15 known as *Myxococcus llanfairpwllgwyngyllgogerychwyrndrobwllllantysiliogogochensis*, or of the crowned
16 slaty flycatcher *Griseotyrannus aurantioatrocristatus*. These mistakes are more likely when dealing with
17 hyper-diverse samples, like plant census (Dauncey et al. 2016, Wagner 2016, Conti et al. 2021). In addition
18 to binomial names, the same species can be known by many vernacular (common) names, which are
19 language or even region-specific: *Ovis aries*, for example, has valid English vernaculars including lamb,
20 sheep, wild sheep, and domestic sheep.

21 All these considerations are actually important when matching species names both within and across
22 datasets. Let us consider the following species survey of individual fishes, European chub, *Cyprinus*
23 *cephalus*, *Leuciscus cephalus*, *Squalius cephalus*: all are the same species (*S. cephalus*), referred to as one of
24 the vernacular (European chub) and two formerly accepted names now classified as synonyms. A
25 cautious estimate of diversity based on the user-supplied names would give $n = 4$ species, when there is in
26 fact only one.

27 A package with the ability to handle the sources of errors outlined above, and especially while provide an
28 authoritative classification, can accelerate the work of consuming large volumes of biodiversity data. For
29 example, this package was used in the process of developing the *CLOVER* database (Gibb et al. 2021) of
30 host-virus associations, by reconciling the names of viruses and mammals from four different sources,

31 where all of the issues described above were present.

32 Overview of functionalities

33 An up-to-date version of the documentation for `NCBITaxonomy.jl` can be found online at
34 <https://docs.ecojulia.org/NCBITaxonomy.jl/stable/>, including examples and a documentation of
35 every method. The package is released under the MIT license. Contributions can be made in the form of
36 issues (bug reports, questions, features suggestions) and pull requests. The package can be checked out
37 and installed anonymously from the central Julia repository:

```
using Pkg

# This line should go in the Julia configuration file - note that the path
# will be created if it doesn't exist, and will be used to store the
# raw taxonomic table
ENV["NCBITAXONOMY_PATH"] = joinpath(homedir(), "data", "NCBITaxonomy.jl")

Pkg.add("NCBITaxonomy") # Downloading the files may take a long time
```

38 The package will download the most recent version of the NCBI taxonomy database, and transform it into
39 a set of Apache Arrow files ready for use. Note that the `NCBITAXONOMY_PATH` can be specified on a per-project
40 basis, and as long as the package is not re-built, the local set of tables downloaded from NCBI will not
41 change; this way, users can re-run an analysis with a guarantee that the underlying taxonomic backbone
42 has not changed.

43 Improved name matching

44 Name finding is primarily done through the `taxon` function, which admits either a unique NCBI identifier
45 (e.g. `taxon(36219)` for the bogue *Boops boops*), a string (`taxon("Boops boops")`), or a data frame with a
46 restricted list of names (see the next section). The `taxon` method has additional arguments to perform
47 fuzzy matching in order to catch possible typos (`taxon("Boops bops"; strict=false)`), to perform a

48 lowercase search (useful when alphanumeric codes are part of the taxon name, like for some viruses), and
49 to restrict the the search to a specific taxonomic rank.

50 The lowercase search can be a preferable alternative to fuzzy string matching. Consider the string
51 Adeno-associated virus 3b - it has three names with equal distance (under the Levenstein string
52 distance function):

```
julia> similarnames("Adeno-associated virus 3b"; threshold=0.95)
3-element Vector{Pair{NCBITaxon, Float64}}:
  Adeno-associated virus - 3 (ncbi:46350) => 0.96
  Adeno-associated virus 3B (ncbi:68742) => 0.96
  Adeno-associated virus 3A (ncbi:1406223) => 0.96
```

53 Depending on the operating system, either of these three names can be returned; compare to the output of
54 a case insensitive name search:

```
julia> taxon("Adeno-associated virus 3b"; casesensitive=false)
Adeno-associated virus 3B (ncbi:68742)
```

55 This returns the correct name.

56 Name matching output and error handling

57 The taxon function will either return a NCBITaxon object (made of a name and id), or throw either a
58 NameHasNoDirectMatch (with instructions about how to possible solve it, using the similarnames
59 function), or a NameHasMultipleMatches (listing the possible valid matches, and suggesting to use
60 alternativetaxa to find the correct one). Therefore, the common way to work with the taxon function
61 would be to wrap it in a try/catch statement:

```
try
    taxon(name)
    # Additional operations with the matched name
catch err
```

```

if isa(err, NameHasNoDirectMatch)
    # What to do if no match is found
elseif isa(err, NameHasMultipleMatches)
    # What to do if there are multiple matches
else
    # What to do in case of another error that is not NCBITaxonomy specific
end
end

```

These functions will not demand any user input in the form of key presses (though they can be wrapped in additional code to allow it), as they are intended to run on clusters without supervision. The `taxon` function has good scaling using multiple threads. For convenience in rapidly getting a taxon for demonstration purposes, we also provide a string macro, whereby *e.g.* `ncbi"Procyon lotor"` will return the taxon object for the raccoon.

Name filtering functions

As the full NCBI names table has over 3 million entries at the time of writing, we have provided a number of functions to restrict the scope of names that are searched. These are driven by the NCBI *divisions*. For example `nf = mammalfilter(true)` will return a data frame containing the names of mammals, inclusive of rodents and primates, and can be used with *e.g.* `taxon(nf, "Pan")`. This has the dual advantage of making search faster, but also of avoiding matching on names that are shared by another taxonomic group (which is not an issue with *Pan*, but is an issue with *Io* as mentioned in the introduction).

Note that the use of a restricted list of names can have significant performance consequences: compare, for example, the time taken to return the taxon *Pan* (ID 9596) in the entire database, in all mammals, and in all primates:

Names list	Fuzzy matching	Time (ms)	Allocations	Memory allocated
all	no	23	34	2 KiB
	yes	105	2580	25 MiB
mammalfilter(true)	no	0.55	32	2 KiB

Names list	Fuzzy matching	Time (ms)	Allocations	Memory allocated
primatefilter()	yes	1.9	551	286 KiB
	no	0.15	33	2 KiB
	yes	0.3	92	27 KiB

Clearly, the optimal search strategy is to (i) rely on name filters to ensure that search are conducted within the appropriate NCBI division, and (ii) only rely on fuzzy matching when the strict or lowercase match fails to return a name, as fuzzy matching can result in order of magnitude more run time and memory footprint. These numbers were obtained on a single Intel i7-8665U CPU (@ (1.90GHz). Using "chimpanzees" as the search string (the NCBI recognized vernacular for *Pan*) gave qualitatively similar results, suggesting that there is no performance cost associated with working with synonyms or vernacular input data.

Quality of life functions

In order to facilitate working with names, we provide the `authority` function (gives the full taxonomic authority for a name), `synonyms` (to get alternative valid names), `vernacular` (for English common names), and `rank` (for the taxonomic rank).

Taxonomic lineages navigation

The `children` function will return all nodes that are directly descended from a taxon; the `descendants` function will recursively apply this function to all descendants of these nodes, until only terminal leaves are reached. The `parent` function is an "upwards" equivalent, giving that taxon from which a taxon descends; the `lineage` function chains calls to `parent` until either `taxon(1)` (the taxonomy root) or an arbitrary ancestor is reached.

The `taxonomicdistance` function (and its in-place equivalent, `taxonomicdistance!`, which uses memory-efficient re-allocation if the user needs to change the distance between taxonomic ranks) uses the Shimatani (2001) approach to reconstruct a matrix of distances based on taxonomy, which can serve as a rough proxy when no phylogenies are available.

98 **Acknowledgements:** This work was supported by funding to the Viral Emergence Research Initiative
99 (VERENA) consortium including NSF BII 2021909 and a grant from Institut de Valorisation des Données
100 (IVADO), by the NSERC Discovery Grants and Discovery Acceleration Supplement programs, and by a
101 donation from the Courtois Foundation. Benchmarking of this package on distributed systems was
102 enabled by support provided by Calcul Québec (www.calculquebec.ca) and Compute Canada
103 (www.computecanada.ca). TP wrote the initial code, TP and CJC contributed to API design, and all
104 authors contributed to functionalities and usability testing.

105 **References**

- 106 Bezanson, J. et al. 2017. Julia: A Fresh Approach to Numerical Computing. - SIAM Review 59: 65–98.
- 107 Chamberlain, S. A. and Szöcs, E. 2013. Taxize: Taxonomic search and retrieval in R. - F1000Research 2:
108 191.
- 109 Conti, M. et al. 2021. Match Algorithms for Scientific Names in FlorItaly, the Portal to the Flora of Italy. -
110 Plants 10: 974.
- 111 Dauncey, E. A. et al. 2016. Common mistakes when using plant names and how to avoid them. -
112 European Journal of Integrative Medicine 8: 597–601.
- 113 Federhen, S. 2012. The NCBI taxonomy database. - Nucleic acids research 40: D136–D143.
- 114 Gibb, R. et al. 2021. Data Proliferation, Reconciliation, and Synthesis in Viral Ecology. - BioScience in
115 press.
- 116 Lefkowitz, E. J. et al. 2018. Virus taxonomy: The database of the International Committee on Taxonomy of
117 Viruses (ICTV). - Nucleic Acids Research 46: D708–D717.
- 118 Norman, K. E. A. et al. 2020. Taxadb: A high-performance local taxonomic database interface. - Methods
119 in Ecology and Evolution 11: 1153–1159.
- 120 Schoch, C. L. et al. 2020. NCBI Taxonomy: A comprehensive update on curation, resources and tools. -
121 Database in press.
- 122 Shimatani, K. 2001. On the Measurement of Species Diversity Incorporating Species Differences. - Oikos
123 93: 135–147.

- 124 Wagner, V. 2016. A review of software tools for spell-checking taxon names in vegetation databases. -
125 Journal of Vegetation Science 27: 1323–1327.
- 126 Walker, P. J. et al. 2020. Changes to virus taxonomy and the Statutes ratified by the International
127 Committee on Taxonomy of Viruses (2020). - Archives of Virology 165: 2737–2748.